

Akshay Danthi

San Francisco, CA | (628) 468-2958 | [Portfolio](#) | akshaydanthi19@gmail.com | [LinkedIn](#)

SUMMARY

AI Engineer with hands-on experience building production-grade LLM applications, including agents and advanced RAG pipelines, specializing in GenAI evaluation. Backed by 8+ years of experience across software engineering, data analytics, and GenAI systems development.

SKILLS

LLM & GenAI: RAG pipelines, Agentic Orchestration (LangGraph, ADK, Semantic Kernel), LangChain, LLM tool calling, prompt engineering, context engineering, ReAct, fine-tuning (LoRA/PEFT), MCP (Model Context Protocol), A2A

LLM Evaluation & LLMOps: GenAI evaluation pipelines (golden datasets, failure-mode taxonomy - open coding & axial coding, LLM as judge, synthetic dataset creation), A/B testing & experimentation, observability, Arize, Galileo, PostHog

ML & Data: Machine Learning (classification, regression), scikit-learn, feature engineering, DL, PyTorch, pandas

Retrieval & Search: Vector & semantic search, FAISS, Pinecone, Weaviate, Reranking & chunking strategies, GraphRAG

Backend: Python, TypeScript, React JS, SQL, Docker, GKE, GitHub Actions, Git, AWS, Azure, GCP, Vertex AI, System Design

PROFESSIONAL EXPERIENCE

AI Engineer

Oct 2025 - Present

[TheSprintFit](#)

San Francisco, CA

- Built & optimized a RAG pipeline to extract structured investment signals from unstructured PDF pitchdecks, improving **grounding accuracy** by **34%** and **reducing hallucination rate** by **62%** through **adaptive chunking & retrieval reranking**.
- Built a LangGraph-powered multi-agent enrichment workflow combining LLM inference, web search, scraping, and validation to complete investor profiles, **increasing field accuracy to 87%** and enhancing founder-investor matching.
- Built a GenAI **evaluation** pipeline using **LangFuse** & Arize **Phoenix** to assess RAG performance, tracking retrieval quality (**Recall@K**, **nDCG@K**, **Precision@K**), grounding accuracy, and **hallucination rate**, with golden datasets, **LLM-as-judge** scoring, and structured failure taxonomy via open and axial coding.
- Shipped production-grade GenAI systems using React (frontend), FastAPI (backend), PostgreSQL, and Docker, deployed on AWS for scalable infrastructure and rapid iteration.

AI Engineer

Sep 2024 - Oct 2025

[Majorly](#)

San Francisco, CA

- Engineered a self-serve NL to SQL analytics engine that enabled career advisors to query student engagement data, achieving a **95% query success rate** across 100+ queries post launch.
- Built a **real-time S2S interview simulator** using Deepgram's speech-to-speech API with LLM tool-calling for dynamic question/rubric retrieval; optimized streaming + concurrency to achieve **<1s end-to-end latency** for realism.
- Developed LLM powered features including a resume analyzer, a LangGraph based career quiz for personalized recommendations, and RAG recommenders for programs, minors, and upskilling paths.

Data Scientist

Dec 2023 - Sep 2024

Golden Gate University

San Francisco, CA

- Analyzed multi year enrollment, retention, graduation, and student course datasets using Python, SQL, and Pandas; delivered insights that guided program planning, course scheduling, and retention interventions for senior leadership.
- Built predictive models using scikit-learn for enrollment forecasting, retention risk, and course demand estimation, improving planning accuracy and enabling data informed decision making across schools and academic programs.
- Developed Tableau dashboards for the Vice Provost that visualized enrollment trends, retention flows, and career outcomes; automated report distribution with Python scripts to reduce manual reporting work across departments.

Product Manager

Sep 2021 - Sep 2023

Mastercard

Chennai, India

- Developed analytics funnels and automated workflows that increased issuer onboarding throughput from 10 to 50+ orgs per week, with dashboards and event tracking that exposed drop offs and failure points.
- Improved the issuer onboarding platform by enhancing REST API integrations across upstream and downstream services during a migration to Java microservices, boosting system stability and sustaining 85% CSAT.

Business Analyst

Nov 2019 - Sep 2021

HSBC

Chennai, India

- Analyzed high-volume global payments data using **Oracle SQL**, mapping end-to-end flows and exception paths; supported concurrent digital-transformation projects for UK BACS and Saudi ACH.
- Translated operational needs into structured requirements and test scenarios for automated payment processing aligned to SWIFT/ISO 20022, **reducing production defects by 40%**.
- Designed **dashboards** and UAT workflows to de-risk go-lives and monitor regional KPIs across payments operations.

Software Engineer

Nov 2017 - Nov 2019

Lloyds Banking Group

Hyderabad, India

- Developed and maintained **Java services** and **Oracle PL/SQL** components for high-volume SWIFT ingestion and processing; handled **50k+ transactions/day** with **99.9% straight-through processing**.
- Optimized performance to cut average processing time from **5s to 2s per transaction** using an LRU cache.

EDUCATION AND CREDENTIALS

Master of Science (M.S.) in Business Analytics

September 2023 – September 2024

Golden Gate University

San Francisco, CA

Bachelor of Engineering (B.E.) in Computer Science

September 2013 – September 2017

Anna University

Chennai, India

- [Portfolio site](#) includes several side-projects on RAG, Agentic AI, ML models for Prediction & Classification, Case Studies, etc