

Akshay Danthi

San Francisco, CA | (628) 468-2958 | [Portfolio](#) | akshaydanthi19@gmail.com | [LinkedIn](#)

SUMMARY

Forward Deployed Engineer who bridges product strategy and deep engineering to ship GenAI that actually gets used. 8 years spanning software engineering, data science, and product leadership, translating ambiguous stakeholder needs into scoped bets, success metrics, and production systems. I build and deploy RAG and agentic workflows end to end, then close the loop with evals, observability, and failure mode analysis to improve grounding, reliability, and adoption.

SKILLS

GenAI & LLMs: RAG, agentic workflows, LangGraph, ADK, LangChain, tool and function calling, prompt engineering, context engineering, LLM APIs (OpenAI, Anthropic, Gemini, Llama)

GenAI evals & reliability: GenAI evaluation strategy, LLM as judge, failure mode analysis, observability, A/B testing

Forward deployed delivery: requirements discovery, technical scoping, solution design, proof of concepts, customer demos, implementation, rollout, documentation, rapid prototyping - v0.dev, lovable, n8n, figma make

Retrieval & data: vector and semantic search, hybrid search, FAISS, Pinecone, Weaviate, pgvector, chunking, embedding pipelines, SQL, Postgres, Pandas

Backend & integration: Python, TypeScript, JS, React JS, FastAPI, REST APIs, webhooks, Docker, CI/CD, Git, GitHub Actions

Cloud & tooling: AWS, GCP, Vertex AI, Azure, MLOps and LLMOps, experiment tracking, PostHog, Arize, Galileo, Helicone

PROFESSIONAL EXPERIENCE

AI Engineer TheSprintFit LLC	Oct 2025 - Present San Francisco, CA
- Partnered with founders and analysts to operationalize founder-investor matching workflows, building an AI system that extracts structured investment signals from unstructured pitch decks; improved grounding accuracy 34% and reduced hallucinations 62%. - Architected a LangGraph-based multi-agent pipeline orchestrating LLM reasoning, web search, and scraping to autonomously enrich investor profiles, increasing dataset accuracy to 87% and improving founder-investor matching. - Implemented a production GenAI evaluation and observability framework using Langfuse, with golden datasets, LLM-as-judge scoring, and structured failure taxonomy to monitor retrieval quality, grounding accuracy, and hallucination rates. - Delivered end-to-end AI platform features across React, FastAPI, PostgreSQL, Docker, and AWS, enabling analysts to explore and evaluate investment opportunities through AI-augmented workflows.	
AI Engineer Majorly Inc.	Sep 2024 - Oct 2025 San Francisco, CA
- Partnered with university career services teams, recruiter SMEs to turn vague “career readiness” workflows into concrete GenAI features, scoping requirements, and success metrics for interview prep and career guidance use cases. - Engineered a self-serve NL to SQL analytics engine that enabled career advisors to query student engagement data, achieving a 95% query success rate across 100 plus queries post launch. - Built a real-time S2S interview simulator using Deepgram’s speech-to-speech API with LLM tool-calling for dynamic question/rubric retrieval; optimized streaming + concurrency to achieve <1s end-to-end latency for realism.	
Data Scientist Golden Gate University	Dec 2023 - Sep 2024 San Francisco, CA

- Analyzed multi year enrollment, retention, graduation, and student course datasets using **Python, SQL, and Tableau**; delivered insights that guided program planning, course scheduling, and retention interventions for senior leadership.
- Built predictive models using scikit-learn for enrollment forecasting, retention risk, and course demand estimation, improving planning accuracy and enabling data informed decision making across schools and academic programs.
- Developed Tableau dashboards for the Vice Provost that visualized enrollment trends, retention flows, and career outcomes; automated report distribution with Python scripts to reduce manual reporting work across departments.

Product Manager

Sep 2021 - Sep 2023

Mastercard

Chennai, India

- Developed analytics funnels and automated workflows that increased issuer onboarding throughput from 10 to 50 plus orgs per week, with dashboards and event tracking that exposed drop offs and failure points.
- Improved the issuer onboarding platform by enhancing REST API integrations across upstream and downstream services during a migration to Java microservices, boosting system stability and sustaining 85% CSAT.

Business Analyst

Nov 2019 - Sep 2021

HSBC

Chennai, India

- Analyzed high-volume global payments data using **Oracle SQL**, mapping end-to-end flows and exception paths; supported concurrent digital-transformation projects for UK BACS and Saudi ACH.
- Translated operational needs into structured requirements and test scenarios for automated payment processing aligned to SWIFT/ISO 20022, **reducing production defects by 40%**.
- Designed **dashboards** and UAT workflows to de-risk go-lives and monitor regional KPIs across payments operations.

Software Engineer

Nov 2017 - Nov 2019

Lloyds Banking Group

Hyderabad, India

- Developed and maintained **Java services** and **Oracle PL/SQL** components for high-volume SWIFT ingestion and processing; handled **50k+ transactions/day** with **99.9% straight-through processing**.
- Optimized performance to cut average processing time from **5s to 2s per transaction** using an LRU cache.

EDUCATION AND CREDENTIALS

Master of Science (M.S.) in Business Analytics

September 2023 – September 2024

Golden Gate University

San Francisco, CA

Bachelor of Engineering (B.E.) in Computer Science

September 2013 – September 2017

Anna University

Chennai, India

- [Portfolio site](#) includes several side-projects on RAG, Agentic AI, ML models for Prediction & Classification, Case Studies, etc